

RTOS Documentation

Kathleen Chung

`kk1chung@uwaterloo.ca`

Connor Cimowsky

`ccimowsky@uwaterloo.ca`

Christian De Angelis

`cdeangel@uwaterloo.ca`

Jaclyne Ooi

`jpooi@uwaterloo.ca`

April 2, 2014

Abstract

This document describes a real-time operating system for the Keil MCB1700 Evaluation Board. It examines the design of the operating system, documents user-facing primitives, and provides an analysis of system performance.

Contents

1	Design Description	5
1.1	Operating System Initialization	5
1.1.1	Memory Initialization	5
1.1.2	Process Initialization	5
1.2	Memory Management	6
1.2.1	Heap Data Structure	6
1.2.2	Requesting Memory Blocks	6
1.2.3	Releasing Memory Blocks	6
1.3	Process Management	7
1.3.1	Process Control Structures	7
1.3.2	Releasing the Processor	8
1.3.3	Process Priority	8
1.3.4	Interprocess Communication	10
1.4	Interrupt Processes	14
1.4.1	Timer I-Process	14
1.4.2	UART I-Process	15
1.5	System Processes	16
1.5.1	Null Process	16
1.5.2	KCD Process	17
1.5.3	CRT Process	17
1.6	User Processes	18
1.6.1	Test Processes	18
1.6.2	Wall Clock Process	18
1.6.3	Set Priority Command Process	19
1.6.4	Stress Test Processes	19
1.6.5	Time Profiler Process	19
1.7	Testing	19

2	Timing Analysis	20
2.1	Experimental Design	20
2.2	Time Measurements	21
2.2.1	request_memory_block	21
2.2.2	send_message	21
2.2.3	receive_message	21
3	Lessons Learned	22
3.1	Version Control	22
3.2	Simulation	22
3.3	Documentation	23
3.4	Linked Structures	23
3.5	Data Structure Allocation	24
	Appendices	24
A	Configuration	25
A.1	Parameters	25
A.2	Macros	25
B	Global Variables	26
B.1	Memory	26
B.2	Processes	26
B.3	Timer I-Process	27
B.4	UART I-Process	28
B.5	KCD Process	29
B.6	Wall Clock Process	29

List of Algorithms

1	Requesting Memory Blocks	6
2	Releasing Memory Blocks	7
3	Releasing the Processor	9
4	Context Switching	10
5	Process Priority	11
6	Sending Messages	12
7	Sending Delayed Messages	13
8	Receiving Messages	13
9	Timer I-Process	14
10	UART I-Process	15
11	Null Process	16
12	KCD Process	17
13	CRT Process	18

List of Figures

1.1	The memory heap	6
1.2	A process control queue (lower priorities omitted)	8
1.3	Data Flow of Input Characters	16
1.4	Data Flow of Output Characters	16

Chapter 1

Design Description

1.1 Operating System Initialization

Once the operating system image has been copied to the LPC1768's on-chip ROM, execution can be initiated by pressing the RESET button on the MCB1700 board. The address of the main function is then loaded into the program counter and the CPU begins fetching and executing instructions. Subsequently, the board's hardware components (e.g., LEDs, timers, serial controllers) are initialized.

1.1.1 Memory Initialization

Now that the board's hardware is ready, it is necessary to initialize global data structures such as process control queues and the keyboard command registry. Most importantly, a portion of the on-chip SRAM is divided into fixed-size blocks (see Appendix A) and inserted into the global memory heap.

1.1.2 Process Initialization


Upon completion of memory initialization, this procedure iterates through the process initialization table and performs three tasks. First, it populates the process control block (PCB) for each process. Next, each PCB is placed in the ready queue (unless it is an i-process, since these processes are invoked by interrupt handlers). Finally, a stack frame is allocated for each process and the resulting stack pointer is stored in its PCB. At this point, the first process can be scheduled by invoking `release_processor`.

1.2 Memory Management

1.2.1 Heap Data Structure

The memory heap structure is implemented using a generic linked list. Each node in the list represents a single memory block that can be requested by invoking the `request_memory_block` primitive and released by invoking the `release_memory_block` primitive. Nodes in the memory heap are spaced apart using a predefined block size (see Appendix A).

Figure 1.1: The memory heap



1.2.2 Requesting Memory Blocks

When a process invokes `request_memory_block`, the operating system first checks if the heap is empty. If so, it will block and then preempt the caller. Otherwise, the next available memory block is popped from the heap and returned. When the process is finished with the memory block, it must release it by invoking `release_memory_block`.

Algorithm 1 Requesting Memory Blocks

```
1: procedure K_REQUEST_MEMORY_BLOCK()
2:   while mem_heap is empty do
3:     K_ENQUEUE_BLOCKED_ON_MEMORY_PROCESS(cur_proc)
4:     K_RELEASE_PROCESSOR()
5:   end while
6:   mem_blk ← POP(mem_heap)
7:   return mem_blk
8: end procedure
```

1.2.3 Releasing Memory Blocks

When a process invokes `release_memory_block`, the operating system first ensures that the provided address points to a valid memory block. If so,

the block is pushed onto the heap. If another process is blocked on memory, it will be unblocked and preemption will take place if necessary.

Algorithm 2 Releasing Memory Blocks

```

1: procedure K_RELEASE_MEMORY_BLOCK(mem_blk)
2:   if mem_blk is invalid then
3:     return RTOS_ERR
4:   end if
5:   PUSH(mem_blk, mem_heap)
6:   if blocked_on_memory_queue is not empty then
7:     blocked_proc ← K_DEQUEUE_BLOCKED_ON_MEMORY_PROCESS()
8:     blocked_proc.state ← READY
9:     K_ENQUEUE_READY_PROCESS(blocked_proc)
10:    K_RELEASE_PROCESSOR()
11:  end if
12:  return RTOS_OK
13: end procedure

```


1.3 Process Management

1.3.1 Process Control Structures

Each process in the operating system is modelled by a process control block. This data structure contains the stack pointer, PID, priority, state, and message queue of a given process.

Depending on their state, processes can be placed in one of three process control queues. The first is the ready queue, which contains processes in the NEW and READY states. The second is the blocked-on-memory queue, which contains processes that are in the BLOCKED_ON_MEMORY state. The third is the blocked-on-receive queue, which contains processes that are in the BLOCKED_ON_RECEIVE state.

Figure 1.2: A process control queue (lower priorities omitted)



1.3.2 Releasing the Processor

Since time slicing is not employed by the operating system, it is frequently necessary for a process to relinquish its usage of the processor. This mechanism is provided by the `release_processor` primitive.

When invoked, this procedure uses the scheduler to select the next process for execution. If the priority of the selected process is greater than or equal to that of the currently executing process, a context switch will occur. Otherwise, the caller will resume execution.

1.3.3 Process Priority

As mentioned in the previous section, each process has a priority which is used to enforce correct precedence during scheduling, blocking, and preemption operations.

Processes may retrieve and modify the scheduling priority of themselves or other processes using two primitives: `get_process_priority` and `set_process_priority`. If a process's priority is changed while it resides in a process control queue, it is removed and then reinserted into the queue corresponding to its new priority. The `release_processor` primitive is then invoked to ensure that preemption will occur if necessary.

Algorithm 3 Releasing the Processor

```
1: procedure K_RELEASE_PROCESSOR()
2:   if ready_queue is empty then
3:     return RTOS_OK      ▷ Do nothing if the ready queue is empty
4:   end if
5:   next_proc  $\leftarrow$  K_DEQUEUE_READY_PROCESS() ▷ Invoke the scheduler
6:   if cur_proc.state  $\neq$  BLOCKED then
7:     if next_proc.priority > cur_proc.priority then
8:       return RTOS_OK      ▷ Do nothing if the priority is lower
9:     end if
10:  end if
11:  K_CONTEXT_SWITCH(cur_proc, next_proc)
12:  return RTOS_OK
13: end procedure

14: procedure K_DEQUEUE_READY_PROCESS()
15:   for i  $\leftarrow$  0 to NUM_PRIORITIES do      ▷ Start at highest priority
16:     if ready_queue[i] is not empty then
17:       return DEQUEUE(ready_queue[i])
18:     end if
19:   end for
20:   return NULL
21: end procedure
```

Algorithm 4 Context Switching

```
1: procedure K_CONTEXT_SWITCH(prev_proc, next_proc)
2:   next_state  $\leftarrow$  next_proc.state
3:   if next_state  $\neq$  NEW and next_state  $\neq$  READY then
4:     return  $\triangleright$  Do nothing if next_proc is unable to execute
5:   end if
6:   if prev_proc.state = EXECUTING then
7:     prev_proc.state  $\leftarrow$  READY
8:     K_ENQUEUE_READY_PROCESS(prev_proc)
9:   end if
10:  prev_proc.sp  $\leftarrow$  __GET_MSP()
11:  next_proc.state  $\leftarrow$  EXECUTING
12:  __SET_MSP(next_proc.sp)
13:  if next_state = NEW then
14:    __RTE()  $\triangleright$  For new processes, pop the exception stack frame
15:  end if
16: end procedure
```

1.3.4 Interprocess Communication

Messages are sent between processes using message envelopes. Each process has a ‘mailbox’, implemented as a queue of message envelopes.

Each message is modelled by two parts: a header used by the operating system, and an envelope used by the sender and recipient. The header contains the PID of the sender and recipient, as well as an expiry time for delayed messages. The envelope contains the message type and the data to be sent.

To send a message, a process must first invoke `request_memory_block` for an envelope. It then populates the envelope and invokes `send_message`, which populates the header and dispatches the message to the recipient. If necessary, the recipient will be unblocked and `release_processor` will be invoked so that preemption may occur; otherwise, control will be returned to the caller. Note that the recipient of the message is responsible for releasing the message envelope by invoking `release_memory_block`.

The procedure for delayed message sending is similar, with the added re-

Algorithm 5 Process Priority

```
1: procedure K_GET_PROCESS_PRIORITY(proc)
2:   if proc.pid is invalid then
3:     return RTOS_ERR
4:   end if
5:   return proc.priority
6: end procedure

7: procedure K_SET_PROCESS_PRIORITY(proc, priority)
8:   if proc.pid is invalid or priority is invalid then
9:     return RTOS_ERR
10:  end if
11:  if proc.priority = priority then
12:    return RTOS_OK      ▷ Do nothing if the priority is unchanged
13:  end if
14:  if proc.state = NEW or proc.state = READY then
15:    REMOVE_FROM_QUEUE(proc, ready_queue)
16:    proc.priority ← priority
17:    K_ENQUEUE_READY_PROCESS(proc)
18:  else if proc.state = BLOCKED_ON_MEMORY then
19:    REMOVE_FROM_QUEUE(proc, blocked_on_memory_queue)
20:    proc.priority ← priority
21:    K_ENQUEUE_BLOCKED_ON_MEMORY_PROCESS(proc)
22:  else if proc.state = BLOCKED_ON_RECEIVE then
23:    REMOVE_FROM_QUEUE(proc, blocked_on_receive_queue)
24:    proc.priority ← priority
25:    K_ENQUEUE_BLOCKED_ON_RECEIVE_PROCESS(proc)
26:  else
27:    proc.priority ← priority
28:  end if
29:  K_RELEASE_PROCESSOR()
30:  return RTOS_OK
31: end procedure
```

Algorithm 6 Sending Messages

```
1: procedure K_SEND_MESSAGE(recipient, msg)
2:   if recipient.pid is invalid then
3:     return RTOS_ERR
4:   end if
5:   if K_SEND_MESSAGE_HELPER(cur_proc, recipient, msg) = 1 then
6:     if recipient.priority  $\leq$  cur_proc.priority then
7:       return K_RELEASE_PROCESSOR()
8:     end if
9:   end if
10:  return RTOS_OK
11: end procedure

12: procedure K_SEND_MESSAGE_HELPER(sender, recipient, msg)
13:  msg.sender  $\leftarrow$  sender
14:  msg.recipient  $\leftarrow$  recipient
15:  ENQUEUE(msg, recipient.msg_queue)
16:  if recipient.state = BLOCKED_ON_RECEIVE then
17:    REMOVE_FROM_QUEUE(recipient, blocked_on_receive_queue)
18:    recipient.state  $\leftarrow$  READY
19:    K_ENQUEUE_READY_PROCESS(recipient)
20:    return 1  $\triangleright$  1 indicates that the recipient was unblocked
21:  else
22:    return 0
23:  end if
24: end procedure
```

quirement that an expiration time is included in the header. When invoked, `delayed_send` places the specified message in the timer i-process's message queue. As described in Section 1.4.1, the timer i-process will dispatch the message at the correct time.

Algorithm 7 Sending Delayed Messages

```

1: procedure K_DELAYED_SEND(recipient, msg, delay)
2:   if recipient.pid is invalid then
3:     return RTOS_ERR
4:   end if
5:   msg.expiry  $\leftarrow$  cur_time + delay
6:   msg.sender  $\leftarrow$  cur_proc
7:   msg.recipient  $\leftarrow$  recipient
8:   ENQUEUE(msg, timer_i_proc.msg_queue)
9:   return RTOS_OK
10: end procedure

```

When a process invokes `receive_message`, the next message will be removed from its message queue and returned. If its message queue is empty, the process is blocked and then preempted using `release_processor`.

Algorithm 8 Receiving Messages

```

1: procedure K_RECEIVE_MESSAGE(sender)
2:   while cur_proc.msg_queue is empty do
3:     K_ENQUEUE_BLOCKED_ON_RECEIVE_PROCESS(cur_proc)
4:     K_RELEASE_PROCESSOR()
5:   end while
6:   msg  $\leftarrow$  DEQUEUE(cur_proc.msg_queue)
7:   sender  $\leftarrow$  msg.sender
8:   return msg
9: end procedure

```

1.4 Interrupt Processes

Both the timer and UART i-processes are scheduled exclusively by their respective interrupt handlers; they are never placed in process control structures (e.g., the ready queue). When an interrupt is received, the following operations are performed:

- The context of the currently executing process is pushed onto the stack
- The appropriate i-process is invoked
- If necessary, the currently executing process is preempted
- The previously saved context is popped off of the stack

1.4.1 Timer I-Process

In order to provide timing services to the operating system, one of the LPC1768's on-chip timers was programmed to signal an interrupt once every millisecond. As described above, this means that the timer i-process will be scheduled at the same interval. The timer i-process is responsible for dispatching delayed messages (sent using `delayed_send`) at the correct time. This is achieved through the use of a time counter and message queues (see appendices B.3.1 and B.3.2).

Algorithm 9 Timer I-Process

```
1: procedure TIMER_I_PROC()
2:    $msg \leftarrow \text{K\_NON\_BLOCKING\_RECEIVE\_MESSAGE}()$ 
3:    $\text{SORTED\_ENQUEUE}(msg, \text{timeout\_queue})$ 
4:   while  $\text{timeout\_queue}$  contains expired messages do
5:      $\text{expired\_msg} \leftarrow \text{DEQUEUE}(\text{timeout\_queue})$ 
6:      $\text{K\_SEND\_MESSAGE\_HELPER}(\text{expired\_msg})$   $\triangleright$  Non-preemptive
7:     if  $\text{expired\_msg.recipient.priority} \leq \text{cur\_proc.priority}$  then
8:        $\triangleright$  preempt  $\text{cur\_proc}$  on completion
9:     end if
10:  end while
11: end procedure
```

1.4.2 UART I-Process

The UART i-process handles interrupts representing two scenarios:

- A character has been entered (see Figure 1.3)
- A character is ready for display (see Figure 1.4)

As such, the UART i-process acts as an interface between the user-facing console and processes in the operating system. To handle input and output, the UART i-process maintains its state between interrupts using global variables (see Appendix B.4).

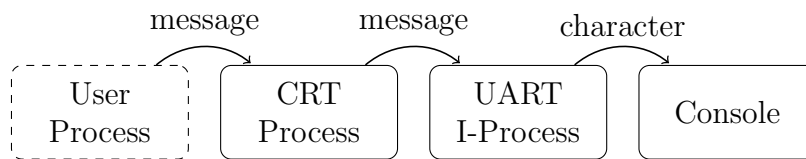
Algorithm 10 UART I-Process

```
1: procedure UART_I_PROC()
2:   if input_char is available then
3:     if mem_heap is not empty then                                ▷ Avoid blocking
4:       msg  $\leftarrow$  K_REQUEST_MEMORY_BLOCK()
5:       msg.type  $\leftarrow$  MSG_TYPE_CRT_DISP
6:       msg.data  $\leftarrow$  input_char
7:       K_SEND_MESSAGE_HELPER(uart_i_proc, crt_proc, msg)
8:       ▷ preempt cur_proc on completion
9:     end if
10:    if input_char  $\neq$  carriage return then
11:      add input_char to input_buffer
12:    else
13:      msg  $\leftarrow$  K_REQUEST_MEMORY_BLOCK()
14:      msg.type  $\leftarrow$  MSG_TYPE_DEFAULT
15:      msg.data  $\leftarrow$  input_buffer
16:      K_SEND_MESSAGE_HELPER(uart_i_proc, kcd_proc, msg)
17:      ▷ preempt cur_proc on completion
18:    end if
19:    else if output_msg is available then
20:      UART0  $\leftarrow$  output_msg.data[output_msg_index]
21:      output_msg_index  $\leftarrow$  output_msg_index + 1
22:    end if
23: end procedure
```

Figure 1.3: Data Flow of Input Characters



Figure 1.4: Data Flow of Output Characters



1.5 System Processes

1.5.1 Null Process

The null process has the lowest priority of any process in the operating system. Since the processor must always be busy, the null process acts as a fail-safe for situations when no other process can be executed. As such, the null process simply invokes `release_processor` in an infinite loop.

Algorithm 11 Null Process

```

1: procedure NULL_PROC()
2:   while true do
3:     RELEASE_PROCESSOR()
4:   end while
5: end procedure

```

1.5.2 KCD Process

The Keyboard Command Decoder (KCD) process provides a console-like interface to users of the operating system. Upon receipt of a command registration message, it uses a global registry to associate the specified command with the message sender. Each time a line of input is entered into the console, the KCD process determines if it is prefixed with a registered command. If so, the input is forwarded to the corresponding process.

Algorithm 12 KCD Process

```
1: procedure KCD_PROC()
2:   while true do
3:      $msg \leftarrow \text{RECEIVE\_MESSAGE}(sender)$ 
4:     if  $msg.type = \text{MSG\_TYPE\_KCD\_REG}$  then
5:        $reg \leftarrow \text{next unused } kcd\_reg \text{ entry}$ 
6:        $reg.id \leftarrow msg.data$ 
7:        $reg.pid \leftarrow sender$ 
8:     else if  $msg.type = \text{MSG\_TYPE\_DEFAULT}$  then
9:        $id \leftarrow \text{first token of } msg.data$ 
10:      if  $kcd\_reg$  contains  $id$  then
11:         $dispatch\_msg \leftarrow \text{REQUEST\_MEMORY\_BLOCK}()$ 
12:         $dispatch\_msg.type \leftarrow \text{MSG\_TYPE\_KCD\_DISPATCH}$ 
13:         $dispatch\_msg.data \leftarrow msg.data$ 
14:         $\text{SEND\_MESSAGE}(kcd\_reg[id].pid, dispatch\_msg)$ 
15:      end if
16:    end if
17:     $\text{RELEASE\_MEMORY\_BLOCK}(msg)$ 
18:  end while
19: end procedure
```

1.5.3 CRT Process

The CRT process forwards text to the system console. In order to achieve this, it repeatedly forwards received messages to the UART i-process. It then triggers a UART output interrupt so that the UART i-process may execute. The mechanism by which the UART i-process achieves interrupt-driven out-

put is outlined in Section 1.4.2.

Algorithm 13 CRT Process

```

1: procedure CRT_PROC()
2:   while true do
3:      $msg \leftarrow \text{RECEIVE\_MESSAGE}()$ 
4:     if  $msg.type = \text{MSG\_TYPE\_CRT\_DISP}$  then
5:        $\text{SEND\_MESSAGE}(uart\_i\_proc, msg)$ 
6:     else
7:        $\text{RELEASE\_MEMORY\_BLOCK}(msg)$ 
8:     end if
9:   end while
10: end procedure

```

1.6 User Processes

1.6.1 Test Processes

In order to ensure correct behaviour of the operating system, six user-level test processes are used to invoke kernel primitives and check for any incorrect results. These test processes use global variables to coordinate with each other and ensure that the operating system works as expected. Some of the key mechanisms that are tested include preemption, modification of process priority, memory block assignment, interprocess communication, and system processes such as the KCD and CRT processes.

1.6.2 Wall Clock Process

The wall clock process uses `delayed_send` to display a digital clock that updates every second. The process sends itself messages with a delay of one second, triggering “tick” updates. Each time the clock ticks, a message is sent to the CRT process using `send_message` to display the current wall clock time. Upon startup, the wall clock process registers three keyboard commands (`%WR`, `%WS`, `%WT`) with the KCD process which allow the time to be set, reset, and stopped.

1.6.3 Set Priority Command Process

Upon startup, this process registers a keyboard command (%C) with the KCD process. This command is handled by invoking `set_process_priority` using the specified process identifier and priority.

1.6.4 Stress Test Processes

When the %Z command is entered, these processes test the behaviour of the operating system under severe memory constraints. If the operating system is implemented correctly, the message “Process C” will be printed to the console every 10 seconds. In order to avoid a deadlock, the priority of process C must be greater than or equal to the priority of process A, as it is responsible for periodically releasing memory blocks.

1.6.5 Time Profiler Process

As described in Section 2.1, this process measures the execution time of `request_memory_block`, `send_message`, and `receive_message`. It can be invoked by entering the %P command.

1.7 Testing

Alongside each deliverable, we implemented a series of automated test processes to evaluate the correctness of our operating system. As described in Section 1.6.1, they test user-facing primitives, system processes, and interrupt processes. Although these tests are very thorough, they are unable to fully assess the interactive components of our operating system, such as the wall clock and debug hotkeys. Therefore, a small portion of our testing was done manually.

When our tests exposed bugs in our operating system, we used the μ Vision debugger to step through our code and track down programming errors. This procedure played a significant role in almost every bug fix, especially those for memory-related bugs (see Section 3.2 for more on this topic).

Chapter 2

Timing Analysis

2.1 Experimental Design

To analyze the performance of our operating system, we programmed the LPC1768's second timer at a frequency of 100 MHz to measure the execution time of three primitives. These primitives are invoked by the time profiler process, which can be initiated using the %P command.

The following code fragment demonstrates the procedure for profiling the execution time of `request_memory_block`:

```
LPC_TIM1->TCR = TIMER_RESET;
LPC_TIM1->TCR = TIMER_START;

for (i = 0; i < NUM_CALLS; i++) {
    request_memory_block();
}

LPC_TIM1->TCR = TIMER_STOP;

/* execution time = LPC_TIM1->TC / NUM_CALLS */
```

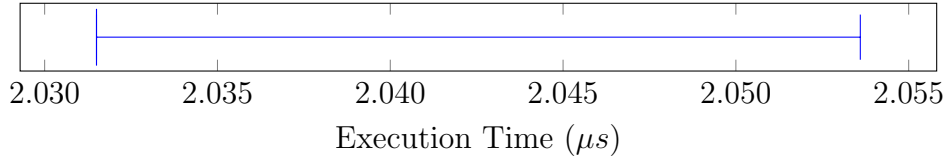
For our analysis, we set `NUM_CALLS` to 100 for each trial; this results in per-call averages with 0.1 nanoseconds of precision. To eliminate the possibility of blocking (and thus preemption), we ensure that `NUM_CALLS` memory blocks are available at the time of execution. For the same reason, we always invoke `send_message` before `receive_message`.

2.2 Time Measurements

The execution time of our trials was very consistent for each primitive, as the Cortex-M3 processor does not employ cache memory. In fact, the processor is described by ARM as “highly deterministic”. Any variance in execution time is most likely due to pipeline stalls (e.g., resulting from branch hazards).

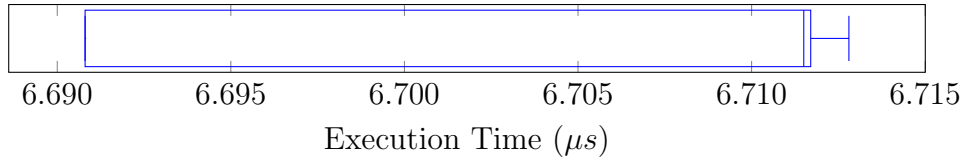
2.2.1 `request_memory_block`

Minimum	First Quartile	Median	Third Quartile	Maximum
2.0315 μs	2.0315 μs	2.0315 μs	2.0315 μs	2.0536 μs



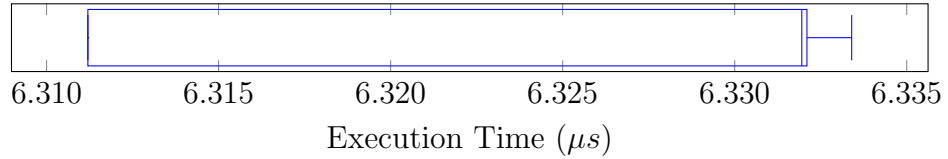
2.2.2 `send_message`

Minimum	First Quartile	Median	Third Quartile	Maximum
6.6908 μs	6.6908 μs	6.7115 μs	6.7117 μs	6.7128 μs



2.2.3 `receive_message`

Minimum	First Quartile	Median	Third Quartile	Maximum
6.3112 μs	6.3112 μs	6.33195 μs	6.3321 μs	6.3334 μs



Chapter 3

Lessons Learned

3.1 Version Control

At the very beginning of our project, we decided to use Git to manage the source code of our operating system. Paired with GitHub, it provided a robust system for ensuring code quality and minimizing bugs. Throughout the project, we obeyed a strict code review policy; all work was done in individual branches, submitted as a pull request, and then exhaustively reviewed by every member of the team before being merged into the `master` branch. This policy allowed us to catch several bugs that would have been incredibly difficult to catch in testing.

3.2 Simulation

For the majority of our first deliverable, we tested our operating system exclusively in the simulator provided by the Keil μ Vision IDE. Before submitting the deliverable, we tested our code on the MCB1700 board; everything worked correctly. This substantiated our assumption that the simulator provided a close approximation to the hardware. In reality, as we discovered during our work on the second deliverable, there is a major difference: the SRAM on the MCB1700 board is not initialized the way it is in the simulator. This led to a number of serious bugs that only surfaced when we tested our code on the board. One of the most common bugs involved dealing with strings. Since we relied on the null character to delimit strings, there were portions of code which would work perfectly in the simulator and not on the

board. These types of bugs only appeared in the second deliverable, since message passing and console I/O both rely heavily on strings. After dealing with these bugs, we learned to write code more defensively and test our code primarily on the board instead of using the simulator.

3.3 Documentation

Beginning with the first deliverable, we decided to write a portion of our documentation alongside each submission. In the time between the first and third deliverables, we compiled a structural description of each part of our project, wrote pseudocode for each major procedure, and completed most of the Lessons Learned section. This work greatly reduced the amount of time it took to complete our final report, as our documentation was nearly complete by the time we submitted the third deliverable.

3.4 Linked Structures

Early on, we knew that we would need to store data in linked structures for different use cases. Since we wanted to avoid coupling these data structures with the type of data they were storing, we chose to make a generic node structure, `node_t`:

```
typedef struct node_t {
    struct node_t *mp_next;
    U32 m_val;
} node_t;
```

We then designed our linked structures to point to nodes of this type. For example, `queue_t`:

```
typedef struct queue_t {
    node_t *mp_first;
    node_t *mp_last;
} queue_t;
```

This way, we can declare new node types with additional members for different use cases. For example, here is a process control block:

```

typedef struct k_pcb_t {
    struct k_pcb_t *mp_next;
    U32 *mp_sp;
    U32 m_pid;
    PRIORITY_E m_priority;
    PROC_STATE_E m_state;
    queue_t m_msg_queue;
} k_pcb_t;

```

As long as the first member is still a pointer to the next node, our linked structures will accommodate these nodes. Creating these generic linked structures saved us a great deal of time in later deliverables.

3.5 Data Structure Allocation

Until the last deliverable, we allocated data structures in the `memory_init` routine. For example, here is the ready queue:

```

queue_t *gp_ready_queue[NUM_PRIORITIES];

for (i = 0; i < NUM_PRIORITIES; i++) {
    gp_ready_queue[i] = (queue_t *)p_end;
    queue_init(gp_ready_queue[i]);
    p_end += sizeof(queue_t);
}

```

This approach unnecessarily adds complexity to our code and also led to quite a few time-consuming bugs at the beginning of the project. To fix it, we began to statically allocate our data structures in the operating system image. Here is the same example using static allocation:

```

queue_t g_ready_queue[NUM_PRIORITIES];

for (i = 0; i < NUM_PRIORITIES; i++) {
    queue_init(&g_ready_queue[i]);
}

```

As a result, our data structures are less error-prone and our code is easier to read. This is the solution that we should have used at the beginning.

Appendix A

Configuration

A.1 Parameters

The following parameters can be adjusted to tune the operating system:

NUM_BLOCKS	Number of memory blocks to be placed in the heap.
BLOCK_SIZE	Size, in bytes, of each memory block in the heap.
USR_SZ_STACK	Size, in bytes, of the stack frame for each process.
NUM_KCD_REG	Maximum number of keyboard commands that can be registered using the KCD process.
KCD_REG_LENGTH	Maximum length of each command identifier.
MSG_LOG_SIZE	Number of recent messages to be logged.
MSG_LOG_LEN	Number of characters to be logged for each message.
NUM_TEST_PROCS	Number of user test processes.

A.2 Macros

The following macros can be defined to achieve additional behaviour:

DEBUG_0	Allow user test processes to print to the console.
DEBUG_1	Enable verbose output for debugging purposes.
DEBUG_LED	Monitor the memory heap using the on-board LEDs.
DEBUG_HOTKEYS	Enable the debug hotkeys (defined in the README).
DEBUG_TIMER	Scale the timer for debugging in the simulator.
DEBUG_DEMO	Leave the user test processes in a blocked state.

Appendix B

Global Variables

B.1 Memory

B.1.1 **g_heap**

During the memory initialization process, part the LPC1768's on-chip SRAM is divided into memory blocks. In order to efficiently serve the needs of processes, the operating system uses a linked list to manage these blocks.

B.1.2 **gp_heap_begin_addr, gp_heap_end_addr**

These variables are set during the memory initialization process and are used for bounds checking in `release_memory_block`.

B.1.3 **gp_stack**

A pointer to the top of the stack. Used during initialization to provide a stack frame for each process.

B.2 Processes

B.2.1 **g_pcbs**

An array of process control blocks for each process in the system. Indexed by PID for constant time access (e.g., from `send_message`).

B.2.2 g_proc_table

An array containing the information required (e.g., PID, priority, entry point) to initialize the process control block for each process in the system.

B.2.3 gp_current_process

A pointer to the PCB of the currently executing process. Used mainly for process control purposes (e.g., release_processor).

B.2.4 g_ready_queue

Implemented as an array of queues (one queue for each priority) containing pointers to process control blocks. Used by release_processor for scheduling processes.

B.2.5 g_blocked_on_memory_queue

An array of queues (one queue for each priority) containing pointers to process control blocks. Used by request_memory_block and release_memory_block for blocking and unblocking processes.

B.2.6 g_blocked_on_receive_queue

An array of queues (one queue for each priority) containing pointers to process control blocks. Used by send_message and receive_message for blocking and unblocking processes.

B.3 Timer I-Process

B.3.1 g_timer_count

The current system time. Used for determining when delayed messages have expired. Incremented each time a timer interrupt is received (1 ms intervals).

B.3.2 g_timeout_queue

A queue of messages that have not yet expired. Sorted by expiry time. Used for dispatching delayed messages at the correct time.

B.3.3 g_timer_preemption_flag

Used to indicate whether the timer interrupt handler should preempt the currently executing process on completion of the timer i-process.

B.4 UART I-Process

B.4.1 g_input_buffer

A buffer containing characters that have been entered by the user since the last carriage return.

B.4.2 g_input_buffer_index

The next available index of the input buffer. Incremented each time a character is appended to the input buffer.

B.4.3 gp_cur_msg

A pointer to the message currently being printed to the console.

B.4.4 g_output_buffer_index

The index of the next character of gp_cur_msg to be printed to the console.

B.4.5 g_uart_preemption_flag

Used to indicate whether the UART interrupt handler should preempt the currently executing process on completion of the UART i-process.

B.5 KCD Process

B.5.1 **g_kcd_reg**

An array containing keyboard commands that have been registered by other processes (e.g., the wall clock process).

B.6 Wall Clock Process

B.6.1 **g_wall_clock_start_time**

The system time at which the wall clock was started. Used for computing the elapsed time for display on the console.

B.6.2 **g_wall_clock_start_time_offset**

The start time specified by the ‘%WS hh:mm:ss’ command. Used for computing the elapsed time for display on the console.

B.6.3 **g_wall_clock_running**

A flag used to indicate if the wall clock is currently running. Enables support for pausing and resuming the wall clock.